


AI Delegation Boundary

Give agents room to act—without giving away the final say.
A clear, exact path from proposal to applied authority.



Guide : AI Delegation Boundary User Guide

Owner : SOLIFAN

Edition : OpenAI WebMCP Challenge · August 2026

Scope : Scope work, challenge boundaries, preserve human judgment, approve exact revisions, and complete safely.

AUTHORITY / WHO DECIDES WHAT

One workspace. Clear authority.

The Agent can propose greater authority. Only a human can grant it.

AGENT	Inspects, proposes, challenges, and re-tests through registered site tools.
HUMAN	Scopes work, resolves challenges, approves the exact revision, and may apply it directly.
BROWSER	Stores workspace state and enforces Guardrails, regression checks, approved state, and final commit.

Controlled authority path



Before Apply

Current revision, exact approval, and deterministic checks must still match.

The browser commits only after exact revalidation.


The Agent may propose, but it cannot approve.
The browser-resident app is the system of record.

OPERATION MAP / END-TO-END FLOW


From candidate boundary to applied authority

Use the upper path for the shortest completion. The lower path is the optional WebMCP application after approval.

PRIMARY · DIRECT IN-PAGE COMPLETION



OPTIONAL · WEBMCP APPLICATION AFTER APPROVAL



Both paths share one exact-state apply coordinator; direct Apply is primary and WebMCP application is optional.

BEFORE YOU BEGIN / WORKING ENVIRONMENT

Open the app and keep one ChatGPT conversation

The browser holds authority; ChatGPT provides the Agent workflow.

VIEW

Open solidfdn.github.io/ai-delegation-boundary-webmcp/ and confirm WebMCP available with 5 TOOLS.

ACT

Open one ChatGPT Desktop conversation—or another WebMCP-compatible Agent client—and keep using it for this workspace.

DONE


Keep the browser tab and that conversation open until the page shows APPLIED.

SOLIFAN | AI Delegation Boundary A pull request for agent authority.

AI Delegation Boundary

A pull request
for agent authority.

Human judgment stays final. Agent proposals are challenged, re-tested, and bound to exact approval.



If WebMCP or 5 TOOLS is missing, stop and confirm the Agent client before scoping work.

01 / HUMAN · TASK SCOPE

Define one task—not an entire function

A clear scope gives every later decision a stable reference.

VIEW

Open Current boundary → Work to delegate.

ACT

Enter one task. Example: “Customer refund decisions.” Add only the policy and constraints that matter.

DONE

Select Start with this work. The five tools remain registered; change, challenge, and review actions can now proceed.

01 HUMAN AUTHORITY **The human scopes the work first.** AGENT TOOLS ARE LOCKED

WHERE 01 · Current boundary → Work to delegate

HUMAN-ONLY ACTION

Current boundary
Review and adjust agent authority

WORK TO DELEGATE

What work or decision are you considering delegating?

Customer refund decisions

Context or constraints (optional)

Decide when an AI agent may complete a refund without human review. Preserve irreversible actions and unknown-policy cases for human control.

Start with this work

Human-only step. Until this is set, the Agent cannot change delegation authority.

CURRENT POLICY
AI may complete only low-impact, reversible, well-supported standard decisions.

Never delegate irreversible execution

02 Review the change
Guardrails, past decisions, and new challenges

Guardrail violations Regressions Challenge gate

WAITING

GUARDRAILS
Non-negotiable boundaries

☐ Irreversible actions remain human-controlled
The agent must never independently execute an irreversible decision.

☐ Unknown policy requires human review
A decision not covered by a known policy cannot be completed autonomously.

KNOWN DECISIONS
Human decisions preserved as tests

No human judgment has been recorded yet.
When a human answers an Agent Challenge, that judgment becomes a regression test for future boundary changes.

AGENT CHALLENGES
Questions that challenge the boundary

First, a human must scope the work.

03 Revision history
Past revisions stay available

v1
Awaiting human task scope
SYSTEM · current

WEBMCP
WebMCP available

Site tools ready

SITE TOOLS
Five normal WebMCP tools are registered by this page. Human Approval can expose one additional apply capability.

AGENT
A compatible Agent client calls the WebMCP tools exposed by this page. No AI backend.

AUTHORITY
The web app owns the boundary, Guardrails, human judgments, approval, and applied state.

NORMAL 5 **AFTER HUMAN APPROVAL** 6

● TASK · CUSTOMER REFUND DECISIONS

Avoid broad scopes such as “Handle all customer service.” Scope is a human-only decision.

02 / HUMAN · STARTING BOUNDARY

Read the starting boundary

It is a candidate for review—not permission to act.

VIEW

Read the Current policy and every boundary rule.

ACT

Confirm both Guardrails: irreversible execution stays human-controlled; unknown policy requires human review.

DONE

You understand the candidate authority for low-risk, reversible, well-supported standard decisions.

00 CURRENT WORKSPACE **Human authority is visible at every layer.** DELIBERATE VERTICAL PAN

Never delegate irreversible execution
BOUNDARY + WEBMCP
Do not delegate

Reversibility = IRREVERSIBLE ▼

Unknown policy requires human review
Human review required

Policy clarity = UNKNOWN ▼

Delegate low-risk standard decisions
Agent may complete

Evidence quality ≥ HIGH ▼

Impact if wrong ≤ LOW ▼

Reversibility = REVERSIBLE ▼

Policy clarity = CLEAR ▼

Exceptionality = STANDARD ▼

Otherwise **Human review required**

First, a human must scope the work.
Until the work is scoped, Agent tools that change, challenge, or review authority are locked.

DECISION PATCH
Every human override becomes a test before it becomes a rule.

Run guardrail & regression checks

The starting boundary is a proposal—not approval.

03 / AGENT · WEBMCP START

Invite ChatGPT to test the boundary

You do not need to type WebMCP tool names.

VIEW


Confirm that the page shows WebMCP available and 5 TOOLS.

ACT

Send: "Help me determine a safe delegation boundary for the work already defined. Use the site tools until human judgment is required."

DONE

ChatGPT inspects, proposes, challenges, and stops for human judgment.



The Agent can recommend and challenge. It cannot scope, judge, or approve.

A proposal is evidence—not permission.

Check the revision number, proposed boundary, Guardrails, regressions, and Challenge gate.

Compare the proposal with the scoped task, consequences, and context.

A concrete Agent Challenge is open for human judgment.



05 / HUMAN · JUDGMENT

Make the judgment the model cannot

Resolve the real-world condition the boundary does not represent.

VIEW

Read the open challenge, its rationale, and modeled factors.

ACT

Record the intended outcome: Allow agent-only (Agent may complete), Keep human review, or Do not delegate.

DONE

The challenge is resolved; your judgment becomes a Known Decision. Authority does not change yet.

07 FRESH EVIDENCE

A new challenge tests the exact revised boundary.

HUMAN DECISION REQUIRED

AI may complete only low-impact, reversible, well-supported standard decisions.

Never delegate irreversible execution

Do not delegate

Reversibility

IRREVERSIBLE

Unknown policy requires human review

Human review required

Policy clarity

UNKNOWN

Otherwise

Human review required

Human decision: Routine verified refund within a concentrated pattern of repeated low-value refunds

Evidence quality: HIGH Exceptionality: STANDARD Impact if wrong: LOW

HUMAN REVIEW REQUIRED

AGENT CHALLENGES

Questions that challenge the boundary

OPEN

Single routine verified low-value refund with no suspicious refund history

After removing the conflicting agent-only rule, this ordinary one-off refund now requires human review even though the evidence is high, impact is low, the action is reversible, policy is clear, and the case is standard. Unlike the preserved Human Decision, there is no concentrated repeated-refund pattern. That context is not represented by the five modeled factors. A human must decide whether blanket review is the intended safe boundary or whether the model needs an explicit aggregate-risk factor before limited agent-only authority can be restored.

Evidence quality: HIGH Exceptionality: STANDARD Impact if wrong: LOW Policy clarity: CLEAR

Reversibility: REVERSIBLE

HUMAN DECISION REQUIRED - Choose one outcome

Allow agent-only

Keep human review

Do not delegate

Run guardrail & regression checks

WEBMCP

WebMCP available

3 TOOLS

Site tools ready

SITE TOOLS

Five normal WebMCP tools are registered by this page. Human Approval can expose one additional apply capability.

AGENT

A compatible Agent client calls the WebMCP tools exposed by this page. No AI backend.

AUTHORITY

The web app owns the boundary. Guardrails, human judgments, approval, and applied state.

NORMAL

5

AFTER HUMAN APPROVAL

6

WHY IT MATTERS

A safe boundary must distinguish ordinary one-offs from aggregate risk—or keep both under review.

You decide the exception; the Agent does not infer your authority.

06 / SYSTEM · KNOWN DECISION

Turn judgment into a durable test

Human decisions protect future revisions.

VIEW

Find the resolved item under Human decisions preserved as tests.

ACT

Confirm that the recorded outcome and decision factors match your intent.

DONE

Future boundary revisions are tested against this Known Decision.

04 HUMAN JUDGMENT

The answer becomes a regression test.

KNOWN DECISION · 1

Unknown policy requires human review

Human review required

Policy clarity = UNKNOWN

Delegate low-risk standard decisions

Agent may complete

Evidence quality ≥ HIGH

Impact if wrong ≤ LOW

Reversibility = REVERSIBLE

Policy clarity = CLEAR

Exceptionality = STANDARD

Otherwise Human review required

FOUNDATIONS CHALLENGES

Questions that challenge the boundary

RESOLVED

Routine verified refund within a concentrated pattern of repeated low-value refunds

All five modeled factors qualify this refund for agent-only completion, while the same customer has received many similar low-value refunds in a short period. The current factor set does not represent cumulative exposure or abuse patterns. A human must decide whether aggregate behavior should override the otherwise qualifying rule.

Evidence quality: HIGH Exceptionality: STANDARD Impact if wrong: LOW Policy clarity: CLEAR

Reversibility: REVERSIBLE

FOUNDATIONS empower challenges.

PRESERVED AS A TEST

Human review now protects future revisions.

Every human override becomes a test before it becomes a rule.

07 / HUMAN + AGENT · RE-TEST

Test the boundary before approval

Passing checks means reviewable—not approved.

VIEW

Return to the same ChatGPT conversation, or locate Run guardrail & regression checks in the page.

ACT

Ask ChatGPT to continue, or run the checks directly. Read **BLOCKED**, **NEEDS_REVIEW**, or **READY_FOR_DECISION**.

DONE

You know whether the revision is safe to consider for approval.

08 RE-TEST

Guardrails pass. Known decisions still hold.

READY FOR DECISION

Human decision: Routine verified refund within a concentrated pattern of repeated low-value refunds

Human review required

Evidence quality: HIGH Exceptionality: STANDARD Impact if wrong: LOW

Human decision: Single routine verified low-value refund with no suspicious refund history

Human review required

Evidence quality: HIGH Exceptionality: STANDARD Impact if wrong: LOW

AGENT CHALLENGES

0

Questions that challenge the boundary

RESOLVED

Single routine verified low-value refund with no suspicious refund history

After removing the conflicting agent-only rule, this ordinary one-off refund now requires human review even though the evidence is high, impact is low, the action is reversible, policy is clear, and the case is standard. Unlike the preserved Human Decision, there is no concentrated repeated-refund pattern. That context is not represented by the five modeled factors. A human must decide whether blanket review is the intended safe boundary or whether the model needs an explicit aggregate-risk factor before limited agent-only authority can be restored.

Agent challenge: Routine verified refund within a concentrated pattern of repeated low-value refunds

AGENT

v1

Human scoped task: Customer refund decisions

Past revision

HUMAN

WebMCP

WebMCP available

5 TOOLS

Site tools ready

SITE TOOLS

Five normal WebMCP tools are registered by this page. Human Approval can expose one additional apply capability.

AGENT

A compatible Agent client calls the WebMCP tools exposed by this page. No AI backend.

AUTHORITY

The web app owns the boundary, Guardrails, human judgments, approval, and applied state.

DETERMINISTIC CHECKS

Guardrail violations

0

Regressions

0

READY_FOR_DECISION means checks passed. It is not approval.

08 / HUMAN + AGENT · RECOVERY

When a revision is blocked

Use only the factors this workspace can represent.

VIEW

Read the generated guidance and check whether an Agent Challenge is still open.

ACT

Resolve any open Challenge first. Otherwise, narrow or remove the conflicting rule using available factors.

DONE

Continue in the same ChatGPT conversation. After any boundary edit, ChatGPT must challenge and review the exact revision again.

05 REGRESSION **The preserved judgment blocks unsafe authority.** BLOCKED · 1 REGRESSION

SOLIFAN | AI Delegation Boundary | A jail request for agent authority. Start new work

AI DELEGATION BOUNDARY

Decide what an AI agent may do on its own.

A human defines the work and retains final authority. The Agent proposes boundary changes, tries to break them with concrete challenges, and re-tests each revision. Only the exact human-approved revision can be applied.

01 HUMAN	02 AGENT	03 HUMAN	04 HUMAN
Scope the work	Propose & Challenge	Decide the boundary	Apply approved revision

01 **Current boundary**
Review and adjust agent authority.

WORK TO DELEGATE

Customer refund decisions

Adapt the conflicting boundary rule

The current boundary conflicts with a human Decision presented from an earlier challenge.

01 - Current boundary → Delegate low-risk standard decisions

02 **Review the change**
Guardrails, past decisions, and new challenges.

Guardrail violations	Regressions	Challenge gate
0	1	PASSED

03 **Revision history**
Past revisions stay available.

v2
Human decided challenge: Routine verified refund within a concentrated pattern of repeated low-value refunds (policy, context)

04 Human Decision.

as a regression test for future revisions:

KNOWN DECISIONS
Human decisions preserved as tests

Human decision: Routine verified refund within a concentrated pattern of repeated low-value refunds Human review

Evidence quality: HIGH · Exceptionality: STRONG · Impact if wrong: LOW

AGENT CHALLENGES
Questions that challenge the boundary

RESOLVED

Routine verified refund within a concentrated pattern of repeated low-value refunds

All five modeled factors qualify this refund for agent-only completion, while the same customer has received many similar low-value refunds in a short period. The current factor set does not reject cumulative exposure or abuse patterns. A human must decide whether aggregate behavior does override the otherwise qualifying rule.

Evidence quality: HIGH · Exceptionality: STRONG · Impact if wrong: LOW · Policy clarity: CLEAR

Reversibility: REVERSIBLE

BOUNDARY CHANGE REFUSED

Guardrails **0**

Known-decision regressions **1**

If safe repair needs a new factor, stop: this workspace cannot model it.

09 / HUMAN · EXACT APPROVAL

Approve only the exact state reviewed

Approval belongs to one revision—not to a general intention.

VIEW

Confirm READY_FOR_DECISION, the final boundary, Known Decisions, and resolved challenges.

ACT

Select Approve revision N.

DONE

Status becomes APPROVED; direct Apply appears; WebMCP changes from five to six tools.

09 EXACT HUMAN APPROVAL

Approval changes the capability surface.

5 — 6 TOOLS

EXACT HUMAN-APPROVED REVISION

v6

This exact revision is human-approved. Only the current approved state can be applied.

Revision history

Past revisions stay available

v6

Human decided challenge: Single routine verified low-value refund with no suspicious refund history

HUMAN · current

Human approval

v5

Agent challenge: Single routine verified low-value refund with no suspicious refund history

AGENT

Past revision

v4

Remove the conflicting agent-only refund rule so the recorded Human Decision remains HUMAN_REVIEW; the available factors cannot safely distinguish its repeated-refund risk.

AGENT

Past revision

v3

Human decided challenge: Routine verified refund within a concentrated pattern of repeated low-value refunds

HUMAN

Past revision

v2

Agent challenge: Routine verified refund within a concentrated pattern of repeated low-value refunds

AGENT

Past revision

v1

Human scoped task: Customer refund decisions

HUMAN

Past revision

WEBMCP

WebMCP available

5 TOOLS

Site tools ready

SITE TOOLS

Five normal WebMCP tools are registered by this page. Human Approval can expose one additional apply capability.

AGENT

A compatible Agent client calls the WebMCP tools exposed by this page. No AI backend.

AUTHORITY

The web app owns the boundary, Guardrails, human judgments, approval, and applied state.

NORMAL

5

→

AFTER HUMAN APPROVAL

6

CONDITIONAL CAPABILITY

apply_approved_revision

Bound to

revision ID + SHA-256 fingerprint

Only the exact approved state can be applied.

Internal SHA-256 binds approval to this exact state and blocks stale application. You never handle the value.

10 / HUMAN · PRIMARY APPLY

Apply in the page—the shortest path

ChatGPT is not required to complete the workflow.

VIEW

Find the card Apply approved revision N.

ACT

Select Apply revision N and complete once.

DONE

The app rechecks the exact approved state, commits APPLIED, and opens the Completion Report.

EXACT HUMAN-APPROVED REVISION**vN**

Only the current approved state can be applied.

PRIMARY COMPLETION PATH**Apply in this page**

The browser rechecks the exact approval before committing.

Apply revision N and complete**● APPLIED · Completion Report**

Direct in-page Apply is the primary completion path.

11 / OPTIONAL · WEBMCP APPLY

Authorize ChatGPT to apply

The Agent may execute only the exact transition you approved.

VIEW

Confirm 6 TOOLS, then use the Send to ChatGPT copy box.

ACT

Send the generated explicit authorization in the same ChatGPT conversation used throughout this workspace.

DONE

ChatGPT calls `apply_approved_revision`; the workspace becomes **APPLIED**; the tool disappears.

10 CHATGPT DESKTOP · WEBMCP

Explicit authority. Exact application.

OPTIONAL AGENT PATH

explicit authorization for this persistent authority change: "continue" was not specific enough. No workaround was attempted.

To proceed, explicitly tell me to apply the approved revision 6.

Human authorization:

Apply the exact currently human-approved revision 6 using the available site tool. Do not modify or substitute the revision. Stop when the workspace shows Applied, or if the approval no longer matches the current revision.

AVAILABLE SITE TOOL

apply_approved_revision

EXACT HUMAN-APPROVED REVISION

v6

This exact revision is human-approved. Only the current approved state can be applied.

CAPABILITY SURFACE

Site tools ready

SITE TOOLS	Five normal WebMCP tools are registered by this page. Human Approval can expose one additional apply capability.
ASERT	A compatible Agent client calls the WebMCP tools exposed by this page. No AI backend.
AUTHORITY	The web app owns the boundary, Guardrails, human judgments, approval, and applied state.
NORMAL	AFTER HUMAN APPROVAL

Explicit authority permits execution—not substitution or re-approval.

12 / SYSTEM · TERMINAL STATE

Know exactly when the work is complete

Completion is a visible terminal state.

VIEW


Find WORKFLOW COMPLETE and COMPLETION REPORT.

ACT

Verify the applied revision, Guardrail violations, protected decisions, resolved challenges, and final boundary.

DONE

The workspace shows APPLIED, WebMCP returns to 5 TOOLS, and no further action is required.



Select Start new work only when you are ready to evaluate a different task.

TROUBLESHOOTING / STATUS AND SCREEN CUES

If the workflow stops, use the screen cue

Each cue tells you what is true now and who should act next.

- | | |
|---|--|
| 01 WEBMCP / 5 TOOLS missing — Stop; use a compatible Agent client and reload. | 05 READY_FOR_DECISION — Checks passed; a human may now approve. |
| 02 WAITING — Scope the work. DRAFT — Follow NEXT for the current step. | 06 APPROVED — Apply directly or explicitly authorize ChatGPT. |
| 03 NEEDS_REVIEW — Follow NEXT; challenge, judgment, or review is still required. | 07 6 TOOLS — Optional WebMCP apply is available after exact approval. |
| 04 BLOCKED — Resolve an open Challenge first; otherwise repair, then re-test. | 08 APPLIED — Review the report; Start new work only for another task. |

Authority should expand only through an exact human decision—and every exception should leave a test behind.